

ASSIGNMENT - 4

| | |
|---------------------|-----------------|
| Assignment Date | 07 OCTOBER 2022 |
| Student Name | SUMITHRA S |
| Student Roll Number | 111519205048 |

Description:- Predicting the age of abalone from physical measurements. The age of abalone is determined by cutting the shell through the cone, staining it, and counting the number of rings through a microscope -- a boring and time-consuming task. Other measurements, which are easier to obtain, are used to predict age. Further information, such as weather patterns and location (hence food availability) may be required to solve the problem.

Attribute Information:

Given is the attribute name, attribute type, measurement unit, and a brief description. The number of rings is the value to predict: either as a continuous value or as a classification problem.

Name / Data Type / Measurement Unit / Description

- 1- Sex / nominal / -- / M, F, and I (infant)
- 2- Length / continuous / mm / Longest shell measurement
- 3- Diameter / continuous / mm / perpendicular to length
- 4- Height / continuous / mm / with meat in shell
- 5- Whole weight / continuous / grams / whole abalone
- 6- Shucked weight / continuous / grams / weight of meat
- 7- Viscera weight / continuous / grams / gut weight (after bleeding)
- 8- Shell weight / continuous / grams / after being dried
- 9- Rings / integer / -- / +1.5 gives the age in years

Building a Regression Model

1. Download the dataset: Dataset
2. Load the dataset into the tool.
3. Perform Below Visualizations.
 - Univariate Analysis
 - Bi-Variate Analysis
 - Multi-Variate Analysis
4. Perform descriptive statistics on the dataset.
5. Check for Missing values and deal with them.
6. Find the outliers and replace them outliers
7. Check for Categorical columns and perform encoding. 8. Split the data into dependent and independent variables. 9. Scale the independent variables
10. Split the data into training and testing
11. Build the Model
12. Train the Model
13. Test the Model
14. Measure the performance using Metrics.

```
#import libraries import numpy as np import pandas as pd
import matplotlib.pyplot as plt import seaborn as sb import
plotly.express as px
```

▼ 2. Load the dataset into the tool

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import os
os.chdir("/content/drive/My Drive")
!ls
```



```
data = pd.read_csv('/content/drive/My Drive/Colab Notebooks/abalone.csv') data
```

[illegible]

4177 rows x 9 columns




▼ 3. Perform Below Visualizations.

- Univariate Analysis

```
data['Rings'].value_counts() data.hist()
```


```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a73e35d0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a6255990>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a620df90>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a61c3bd0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a6188210>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a6140810>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a6173e90>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a6138410>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fa4a6138450>]],  
      dtype=object)
```



Bi-Variate Analysis

```
plt.scatter(data.Rings, data.Sex) plt.title('The
Gender of Abalone vs Number of Rings')
plt.xlabel('No. of Rings') plt.ylabel('Gender')
```


```
Text(0, 0.5, 'Gender')
```



• Multi-Variate Analysis

```
sb.heatmap(data.corr(),annot=True
)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fa4a59c4350>
```



▼ 4. Perform descriptive statistics on the dataset.

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
 # Column          Non-Null Count  Dtype
----  -
0  Sex              4177 non-null  object
1  Length           4177 non-null  float64
2  Diameter         4177 non-null  float64
3  Height           4177 non-null  float64
4  Whole weight     4177 non-null  float64
5  Shucked weight   4177 non-null  float64
6  Viscera weight   4177 non-null  float64
7  Shell weight     4177 non-null  float64
8  Rings            4177 non-null  int64
dtypes: float64(7), int64(1), object(1) memory
usage: 293.8+ KB
```

```
data.describe()
```

| | Diameter | | Height | Whole Shucked | | Viscera | Length |
|-------|-------------|-------------|-------------|---------------|-------------|-------------|-------------|
| | | | | weight | weight | | weight |
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 |
| mean | 0.523992 | 0.407881 | 0.139516 | 0.828742 | 0.359367 | std | |
| | 0.120093 | 0.099240 | 0.041827 | 0.490389 | 0.221963 | min | 0.180594 |
| | 0.075000 | 0.055000 | 0.000000 | 0.002000 | 0.001000 | 25% | 0.109614 |
| | 0.450000 | 0.350000 | 0.115000 | 0.441500 | 0.186000 | | 0.000500 |
| 50% | 0.545000 | 0.425000 | 0.140000 | 0.799500 | 0.336000 | | 0.093500 |
| 75% | 0.615000 | 0.480000 | 0.165000 | 1.153000 | 0.502000 | | 0.171000 |
| | | | | | | | 0.253000 |
| | | | | | | | 0.760000 |
| max | 0.815000 | 0.650000 | 0.650000 | 1.130000 | 0.650000 | | 0.760000 |
| | 2.825500 | | 1.488000 | | | | |

5. Check for Missing values and deal with them.

There is no missing values


```
data.isnull().any()
```

```
Sex           False
Length        False
Diameter       False
Height         False
Whole weight   False
Shucked weight False
Viscera weight False
Shell weight   False
Rings dtype:   False
bool
```

6. Find the outliers and replace them outliers

The dataset does not have a outliers

```
fig = px.histogram(data, x='Whole weight')
fig.show()
```



7. Check for Categorical columns and perform encoding.

There is one Categorical column SEX is replaced by an Integer

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
data["Sex"] = le.fit_transform(data["Sex"])
data["Sex"]
```

```
0      2
1      2
2      0
3      2
4      1
```

```
4172    0
4173    2
4174    2
4175    0
4176    2
```

```
Name: Sex, Length: 4177, dtype: int64
```

▼ 8. Split the data into dependent and independent variables.

```
x=data.iloc[:,0:8].values
y=data.iloc[:,8:9].values
```

```
x
```

```
array([[2.      , 0.455 , 0.365 , ..., 0.2245, 0.101 , 0.15 ],
       [2.      , 0.35  , 0.265 , ..., 0.0995, 0.0485, 0.07 ],
       [0.      , 0.53  , 0.42  , ..., 0.2565, 0.1415, 0.21 ],
       ...,
       [2.      , 0.6   , 0.475 , ..., 0.5255, 0.2875, 0.308 ],
       [0.    , 0.625 , 0.485 , ..., 0.531 , 0.261 , 0.296 ],
       [2.    , 0.71  , 0.555 , ..., 0.9455, 0.3765, 0.495 ]])
```

```
y
```

```
array([[15], [ 7],
       [ 9], ...,
       [ 9],
       [10],
       [12]])
```

▼ 9. Scale the independent variables

```
x=data.iloc[:,0:8]
print(x.head())
```

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight \ |
|---|-----|--------|----------|--------|--------------|------------------|
| 0 | 2 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 |
| 1 | 2 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 |
| 2 | 0 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 |
| 3 | 2 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 |

| | | | | | | |
|---|---|----------------|--------------|-------|--------|--------|
| 4 | 1 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 |
| | | Viscera weight | Shell weight | | | |
| 0 | | 0.1010 | 0.150 | | | |
| 1 | | 0.0485 | 0.070 | | | |
| 2 | | 0.1415 | 0.210 | | | |
| 3 | | 0.1140 | 0.155 | | | |
| 4 | | 0.0395 | 0.055 | | | |

▼ 10. Split the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=0)
```

```
x_train.shape
```

```
(2923, 8)
```

```
x_test.shape
```

```
(1254, 8)
```

▼ 11. Build the Model

```
from sklearn.linear_model import LinearRegression lr =
LinearRegression()
```

▼ 12. Train the Model

```
lr.fit(x_train, y_train)
```

```
LinearRegression()
```

▼ 13. Test the Model

```
y_pred = lr.predict(x_test)
print((y_test)[0:6])
print((y_pred)[0:6])
```



```

[[13]
 [ 8]
 [11]
 [ 5]
 [12]
 [11]]
[[13.11640829]
 [ 9.65691091]
 [10.35350972]
 [ 5.63648715]
 [10.67436485]
 [11.95341338]]

```

▼ 14. Measure the performance using Metrics.

```
# RMSE(Root Mean Square Error)
```

```

from sklearn.metrics import mean_squared_error mse =
mean_squared_error(y_test, y_pred) rmse = np.sqrt(mse)
print("RMSE value : {:.2f}".format(rmse))

```

```
RMSE value : 2.26
```

```

from sklearn.model_selection import cross_val_score
cv_scores = cross_val_score(lr, x, y, cv=5)
sco=cv_scores.round(4) print(cv_scores.round(4))
print("Average",sco.sum()/5)

```

```
[0.4113 0.1574 0.4807 0.5046 0.4362] ----- Average 0.39803999999999995
```